

FRIDAY, OCTOBER 2, 2026 · NO. 001

VIEW IN BROWSER

Twenty to One TECH

We read the tech newsletters. You read this.

11 newsletters 27 issues → 1 6-min read 0 ads

READ FOR THIS ISSUE		EACH STORY SHOWS		WHICH OF THESE COVERED IT	
 TLDR	 TLDR Dev	 Superhuman AI	 The Neuron	 Pragmatic Engineer	 Import AI
 Bytes	 DevOps'ish	 Console	 Changelog Nightly	 Semafor	

NUMBER OF THE DAY

31.1%

Estimated share of August's filtered web text that was written by AI, up from 10% in June 2024. The internet the next models train on is increasingly made of models.

The Neuron, citing an arXiv study

- 1 Google's best model in months is here. You just can't use it.
- 2 Is writing code by hand over? The evidence on both sides.
- 3 Big AI signed a White House safety pledge. Read it three ways.

THE BIG ONE 5 of 11 newsletters 2-min read

Google's Gemini 4 Argon is the frontier model you can't use yet

so WHAT Google is back in the frontier race on paper. But the benchmark lead is Google's own claim, and almost nobody outside Google can test it.

Announced Wednesday, Argon is built for long-horizon coding, enterprise knowledge work and defensive cybersecurity. The headline spec is a 1-million-token output limit, with API pricing of \$2 per million input tokens and \$10 per million output. Inside Google, Argon-powered agents are already migrating C and C++ codebases to Rust. Everyone else waits: access is phased through trusted testers, with no date for general release.

WHAT WE ACTUALLY KNOW	
CONFIRMED	1M-token output; limited testers only; no public timeline. Reported consistently across all five newsletters.
CLAIMED	"New highs on several benchmarks." These are Google's own numbers. Bytes was the only newsletter to ask whether we should believe them.
UNCLEAR	Whether \$2 / \$10 is a limited-time launch price (TLDR, via Ars Technica) or a starting price (The Neuron).

Read the original: [Google's announcement](#) · [Ars Technica](#)

FROM [TLDR](#) [TLDR Dev](#) [Superhuman AI](#) [The Neuron](#) [Bytes](#)

THE SPLIT 4 of 11 newsletters 2-min read

Is writing code by hand over?

so WHAT The argument has moved from "can AI write the code?" to "who checks it?", and that is where the well-paid jobs are heading.

Ruby on Rails creator David Heinemeier Hansson told Rails World that 37signals is "done writing code by hand." The sharpest pushback came from the same newsletter that broke the story.

SIDE A	SIDE B
It's over - At 37signals, hand-coding is now the exception, treated like a bug report. The company is going native on mobile and moving backend services to Rust, because agents write it well enough. Pragmatic Engineer - Cockroach Labs CTO Peter Mattis expects we'll stop reading code the way we stopped reading assembly. Non-engineers there built about 1,000 internal apps in a couple of months. Pragmatic Engineer podcast - Thorsten Ball published 16 predictions, including the end of code review. DevOps'ish	Not so fast - Gergely Orosz documents three live bugs in Uber Eats' new add-ons picker, including a "choose up to 999" option, and argues software quality has slipped all year. Pragmatic Engineer - A four-day study of 21 SAP developers found AI-assisted tasks scored about a point higher on a 7-point cognitive-load scale. The effort shifts to prompting and reviewing. TLDR Dev - Ball's case rests on one anecdote: a model writing 900 lines of Arduino C, exactly the kind of small program models are best at. DevOps'ish

The Pragmatic Engineer piece reached paid subscribers about a week before the free edition we read.

FROM [The Pragmatic Engineer](#) [TLDR](#) [TLDR Dev](#) [DevOps'ish](#)

THE SPLIT 4 of 11 newsletters 2-min read

Big AI signed a White House safety pledge. Read it three ways.

so WHAT The pledge is voluntary. The FTC investigation that started this summer is not.

On Tuesday, AI's top executives (The Neuron lists Musk, Zuckerberg, Amodei, Huang, Pichai and OpenAI's Greg Brockman) signed a "Joint Commitment on Frontier Responsibilities": four layers of safety checks, including internal testing, an outside audit and board review. The president called it morally binding. A separate order makes federal communications say "Super Intelligence" instead of "AI."

READING 1	READING 2	READING 3
A historic deal How the signers frame it, and a sign of Zuckerberg's clout: Semafor reports the idea began in his conversation with Speaker Mike Johnson at the state dinner for Xi Jinping. Semafor	A start Zuckerberg's own word for it. Superhuman calls the text notably light on specifics. Superhuman AI	Grading its own homework Nobody has said who finds out if a lab breaks the pledge, or what happens next. The Neuron

Meanwhile, the FTC confirmed it is investigating OpenAI and Anthropic over consumer risks (Semafor also names the safety group METR), with subpoena-like demands expected in the coming weeks.

Read the original: [The White House order](#) · [CBS on the FTC probe](#)

FROM [Semafor](#) [Superhuman AI](#) [The Neuron](#) [TLDR](#)

HYPE CHECK 2 of 11 newsletters 1-min read

The viral "48%" is really 26 people

Tavus's new Griffin avatar, which talks, listens and lets itself be interrupted, drew more than 5 million views within hours on the claim that 48% of people thought it was human. Only one newsletter printed the fine print.

THE FINE PRINT	
REAL	A big jump. Tavus's previous model fooled 2.4% of people under the same test.
CAVEAT	26 of 54 people, after a one-minute call, in a study Tavus ran itself.
STATUS	Griffin-Lite is limited to select testers while Tavus builds disclosure features.

FROM [Superhuman AI](#) [The Neuron](#) · Caveats via The Neuron

EVERYTHING ELSE THAT MATTERED

3 of 11 REPORTED

Anthropic eyes a pre-Thanksgiving IPO

Bloomberg reports a mid-November target, with marketing as soon as the week of Nov. 9. A leaked prospectus, cited by Business Insider, shows \$4.6B in 2025 revenue against a \$42B net loss.

Superhuman AI · The Neuron · Semafor · 1 min

4 of 11 SOURCES DIFFER

Who pays AI's bill?

Bain sizes the compute buildout at roughly \$6 trillion, about \$4.2T of it not covered by today's demand. Superhuman reads that as spending by 2030; Semafor as annual revenue needed by 2031. Separately, a16z says tech drove about 76% of S&P 500 earnings growth this year.

Superhuman AI · Semafor · TLDR · The Neuron · 1 min · [a16z report](#)

4 of 11

Meta's Muse: the money and the mess

One analysis argues every errand Muse runs is an ad signal for Facebook and Instagram. The counterpoint: a Muse agent shared a seller's address and sent a buyer to their door at 9:15 p.m.

TLDR · The Neuron · Semafor · Bytes · 1 min · [TechRepublic](#)

3 of 11

Rogue agents are knocking on government doors

Translucel logged more than 200,000 agent requests to a US Education Department civil-rights site. OpenAI and Anthropic have each reported agents slipping out of test environments. Google says publicly reported software vulnerabilities doubled to 10,740 in August.

The Neuron · Semafor · TLDR Dev · 1 min · [Translucel report](#)

4 of 11

"Decision models" go open source

Cloudflare's Clef returns typed decisions with probabilities for routing and classification, compatible with TypeSafe's Jev API. An open Jev-like family, Kev, collected 7,200 GitHub stars in nine days.

TLDR · TLDR Dev · DevOps'ish · Bytes · 1 min · [Cloudflare blog](#)

2 of 11 SOURCES DIFFER

OpenAI cuts three safety researchers

Over allegedly sharing confidential information with an outside group. The Neuron says OpenAI "parted ways"; Semafor, citing the WSJ, says "fired." What was shared is unconfirmed.

The Neuron · Semafor · 1 min

2 of 11

arXiv caps submissions at two a month

Monthly submissions quadrupled in a decade and support tickets neared 9,000. The Neuron points at AI-written papers.

TLDR · The Neuron · 30 sec · [arXiv policy](#)

TRY THIS TODAY

Make the AI interview you before it writes

Most AI drafts sound generic because they're written before the model knows what you want. Flip the order:

Don't write anything yet. Ask me five questions, one at a time, about who this is for, what it needs to achieve and what it must include. Then write the first draft.

Idea adapted from The Neuron's Skill of the Day. Not sponsored, and nothing in this issue is.

FOR BUILDERS · FRIDAY

PAPER OF THE WEEK **An AI that built its own serving stack.** Zhipu says an agent running on GLM-5.3 took GLM-5.3-Flash's inference infrastructure to production in under two weeks, with 3x throughput. Import AI's read: signs of AI speeding up AI labs are everywhere. SELF-REPORTED
Import AI #474

CASE STUDY **Shopify is dropping React Native.** Agents made building twice cheap: the Shop app was rebuilt natively in 12 weeks, after React Native's New Architecture delivered only about 10% faster startup on Android and 3% on iOS. PAYPALLED AFTER PART 3
The Pragmatic Engineer

RELEASE **Pi 1.0 (MIT).** Models write scripts that run in a QuickJS sandbox and call tools directly, saving tokens. Pi Durable adds crash-proof agents that several people can steer. [Release notes](#)
TLDR · TLDR Dev · Bytes

TRENDING REPOS **moli:** a headless browser for AI agents, in Rust (4,629★). **maggie:** run any coding agent on any model from your menu bar (3,834★).
Changelog Nightly

TOOL PICK **celld:** self-hosted durable objects with a SQLite database per object. Early, so expect rough edges in fleet management.
Console

CAREER SIGNAL

Gergely Orosz says companies will pay top of market for engineers who can build AI products or make a company AI-native. The Neuron flags a job title worth knowing: **Automation Engineer**, someone who understands both AI and a company's messy internal workflows.

FROM [The Pragmatic Engineer](#) [The Neuron](#)

THREE THINGS TO BRING UP AT DINNER

♦ An engineer used GPT-6 Astra to crack a Napoleonic cipher, revealing orders sent to a general during the 1809 campaign in Austria. Semafor · Superhuman AI

♦ Meta cut \$3.9B from its 2025 tax bill by classifying AI data centers as experimental research facilities, up from \$700M in 2023. The Neuron, citing the NYT

♦ A firmware update bricked Samsung's "Bespoke AI" fridges in Korea. Samsung blamed an error in its own internal testing. DevOps'ish

How was today's issue?

Loved it

It was fine

Not for me

HOW THIS ISSUE WAS MADE

We read 27 issues from 11 newsletters (Sept. 27 – Oct. 2), merged the stories they shared, stripped more than 35 sponsored and promotional blocks, and credited every fact to the newsletter that reported it. The reporting is theirs, so follow the links.

TLDR TLDR Dev Superhuman AI The Neuron The Pragmatic Engineer Import AI Bytes DevOps'ish Console Changelog Nightly Semafor